



POLYTECHNIQUE  
MONTRÉAL  
TECHNOLOGICAL  
UNIVERSITY



## Breaking free from the matrix: A stabilized Navier-Stokes perspective using Lethe deal.II workshop

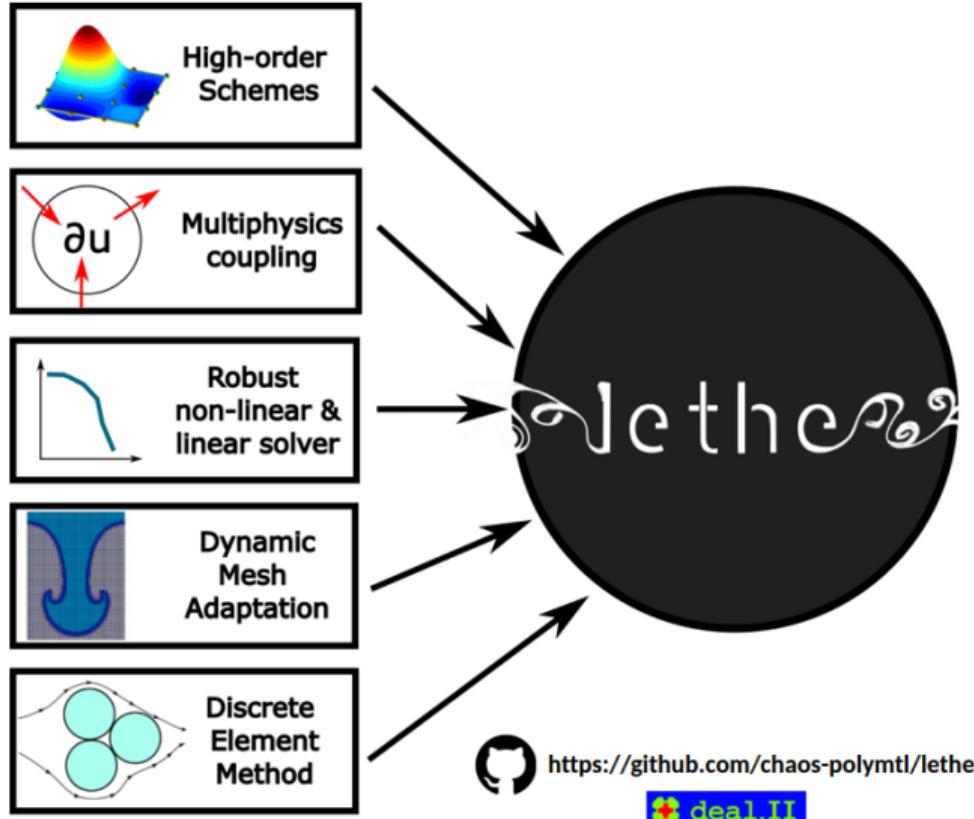
Laura Prieto Saavedra<sup>1</sup> Peter Munch<sup>2</sup> Bruno Blais<sup>1</sup>

<sup>1</sup>Polytechnique Montréal, laura.prieto-saavedra@polymtl.ca, bruno.blais@polymtl.ca

<sup>2</sup>Uppsala University, peter.munch@it.uu.se

Fort Collins, Tuesday 13<sup>th</sup> August, 2024

## Open-source software: Lethe

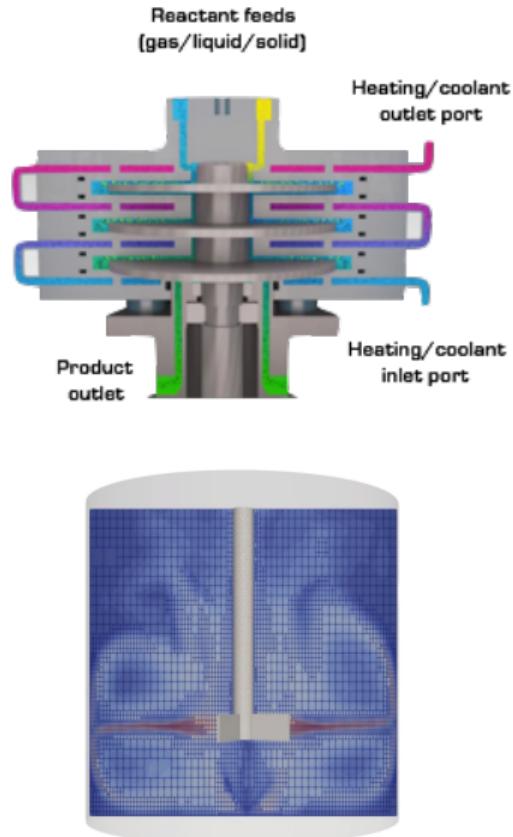


- ▶ Based on the deal.II library (Arndt et al., 2023)
- ▶ Created by Blais et al. (2020)
- ▶ 35 contributors
- ▶ User guide (> 50 examples)
- ▶ Tests (> 350 tests)

## Context and motivation: process-intensified equipment

- ▶ Simulation of complex flow problems in chemical engineering is still a challenge: **rotating flows, transient flows** →  $\text{Re} = [10^3, 10^5]$ .
- ▶ Localized physical phenomena require very fine meshes → alternative: **locally refined meshes**.
- ▶ The algorithms for these very large simulations are memory bandwidth bounded, which led us to implement a **matrix-free solver**.
- ▶ Several challenges arise, among them: the development of a **matrix-free linear solver**.

In this presentation: overview of the matrix-free solver and the geometric multigrid preconditioner



## Outline

1. Governing equations and discretization: incompressible Navier-Stokes equations
2. Implementation: from a matrix-based to a matrix-free approach
3. Linear solver: geometric multigrid preconditioner
4. Numerical tests
5. Outlook: what comes next?

## 1. Governing equations and discretization

Incompressible Navier-Stokes:

$$\nabla \cdot \mathbf{u} = 0 \text{ in } (0, T] \times \Omega$$

$$\partial_t \mathbf{u} + (\mathbf{u} \cdot \nabla) \mathbf{u} + \frac{1}{\rho} \nabla p - \nabla \cdot \mathbb{D}(\mathbf{u}) = \mathbf{f} \text{ in } (0, T] \times \Omega$$

- ▶ Non-linear coupled PDE system that needs to be solved for a vector-valued velocity  $\mathbf{u}$  and a scalar pressure  $p$ .
- ▶ Continuous Galerkin method.
- ▶ Tensor elements: quadrilateral and hexahedra.
- ▶ Newton's method for the non-linearity.
- ▶ For steady-state problems:  $\partial_t \mathbf{u} = 0$ .

## 1. Governing equations and discretization (Cont.)

**Weak form:**

$$\begin{aligned} F(\mathbf{u}, p) := & (q, \nabla \cdot \mathbf{u})_{\Omega} + (\mathbf{v}, \partial_t \mathbf{u})_{\Omega} + (\mathbf{v}, (\mathbf{u} \cdot \nabla) \mathbf{u})_{\Omega} - (\nabla \cdot \mathbf{v}, p)_{\Omega} \\ & + \nu (\nabla \mathbf{v}, \nabla \mathbf{u})_{\Omega} + (\mathbf{n} \cdot \mathbf{v}, p)_{\partial\Omega} - (\mathbf{v} \mathbf{n}, \nu \nabla \mathbf{u})_{\partial\Omega} - (\mathbf{v}, \mathbf{f})_{\Omega} = 0 \end{aligned}$$

**Linear system for each Newton iteration:**

$$\underbrace{\begin{bmatrix} \mathbf{A} & \mathbf{B} \\ \mathbf{B}^T & \mathbf{0} \end{bmatrix}}_{F'} \begin{bmatrix} \mathbf{U} \\ \mathbf{P} \end{bmatrix} = - \underbrace{\begin{bmatrix} \mathbf{R}_u \\ \mathbf{R}_p \end{bmatrix}}_R$$

This discretization has three limitations:

- ▶ Ladyzhenskaya–Babuška–Brezzi condition.
- ▶ Instabilities when Re increases.
- ▶ Saddle-point problem.

→ we overcome these by using stabilization techniques

## 1. Governing equations and discretization (Cont.)

$$F + \underbrace{\sum_k (\tau \nabla \cdot q, R_m)_{\Omega_k}}_{\text{PSPG term}} + \underbrace{\sum_k (\tau (\mathbf{u} \cdot \nabla) \mathbf{v}, R_m)_{\Omega_k}}_{\text{SUPG term}} = 0 \quad R_m : \text{momentum strong residual}$$

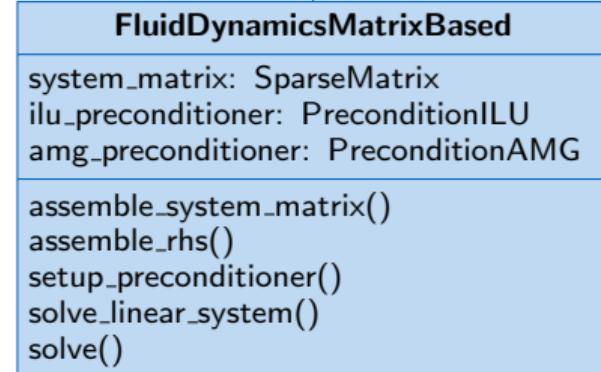
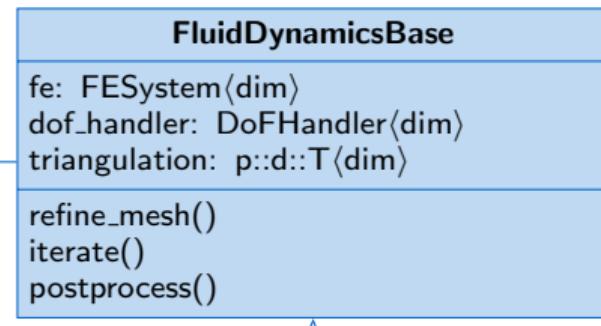
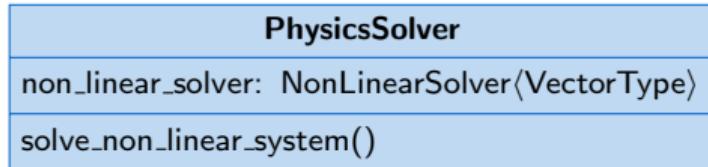
$$\tau = \left[ \left( \frac{1}{\Delta t} \right)^2 + \left( \frac{2\|\mathbf{u}\|}{h} \right)^2 + 9 \left( \frac{4\nu}{h^2} \right)^2 \right]^{-1/2} \quad h = \left( \frac{6h_k}{\pi} \right)^{\frac{1}{3}}$$

Modified linear system:

$$\underbrace{\begin{bmatrix} \hat{\mathbf{A}} & \hat{\mathbf{B}} \\ \mathbf{C} & \mathbf{D} \end{bmatrix}}_{F'} \begin{bmatrix} \mathbf{U} \\ \mathbf{P} \end{bmatrix} = - \underbrace{\begin{bmatrix} \hat{\mathbf{R}}_u \\ \hat{\mathbf{R}}_p \end{bmatrix}}_R$$

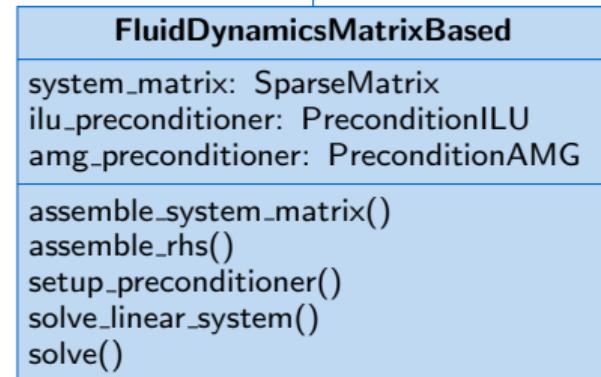
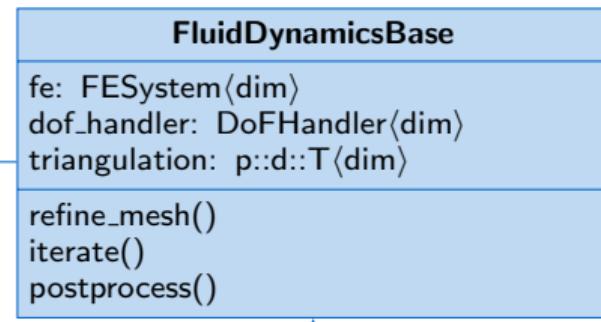
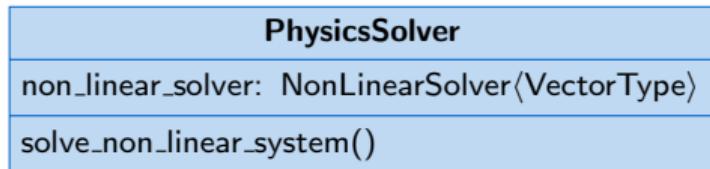
- ▶ We solve this problem in a fully-coupled monolithic way.
  - ▶ At each nonlinear iteration, we solve the system using a preconditioned GMRES method.
- how did we implement such a solver?

## 2. Implementation: from a matrix-based to a matrix-free approach



- ✓ Tempered by dim and VectorType.
- ✓ Physics solver allows us to experiment with different non linear solvers (Newton, inexact, kinsol) and allow us to solve other physics.
- ✓ The solve() function is where everything happens: read/create mesh, setup DoFs, initial conditions, iterate, postprocess.
- ✓ AMG and ILU preconditioners available through TrilinosWrappers.

## 2. Implementation: from a matrix-based to a matrix-free approach



**Note:** We are interested in  $p = 2, 3$  elements.

**Two main issues:**

- ✗ Assembly of the matrix is time-consuming (time increases with the order of FE).
- ✗ Large matrices lead to a memory-bandwidth bottleneck.

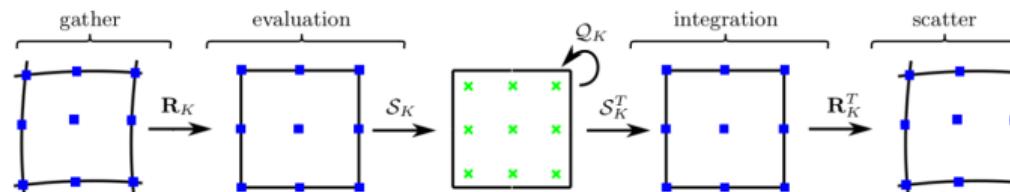
→ how to improve this?

## 2. Implementation: from a matrix-based to a matrix-free approach (Cont.)

**Main idea:** compute results of matrix-vector multiplication on-the-fly by defining an operator.

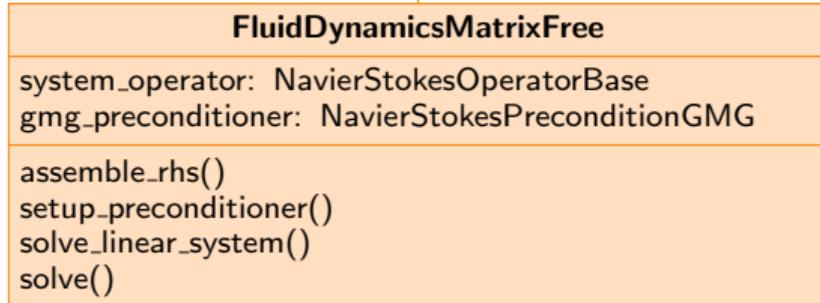
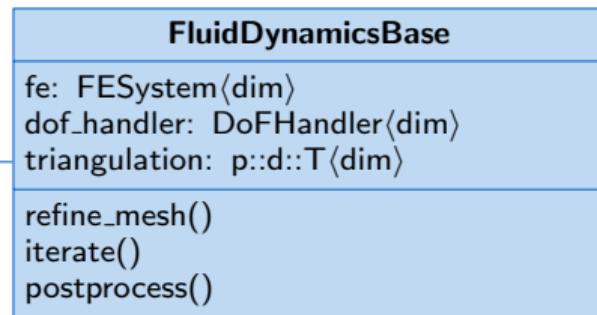
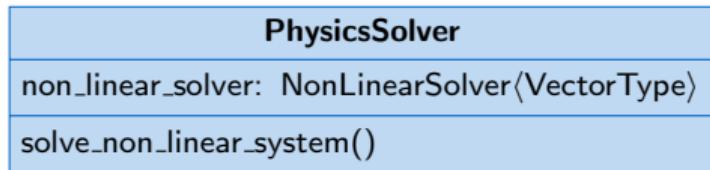
The operator  $\mathcal{A}$  allows us to rewrite the matrix-vector multiplication  $\mathbf{A}\mathbf{u}$  as follows:

$$\mathbf{v} = \mathcal{A}(\mathbf{u}) = \left( \sum_k \mathbf{R}_k^T \mathbf{A}_k \mathbf{R}_k \right) \mathbf{u} = \sum_k \mathbf{R}_k^T \underbrace{\mathbf{S}_k^T \mathbf{Q}_k \mathbf{S}_k}_{\mathbf{A}_k} \mathbf{R}_k \mathbf{u}$$



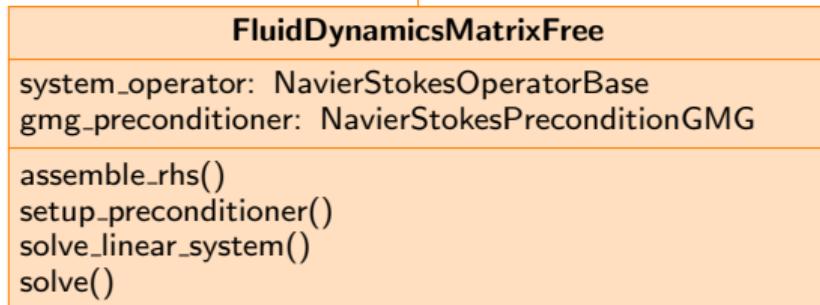
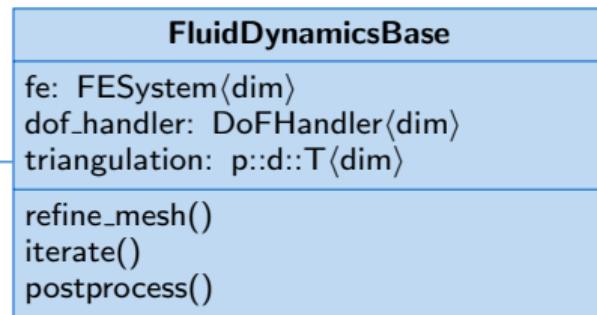
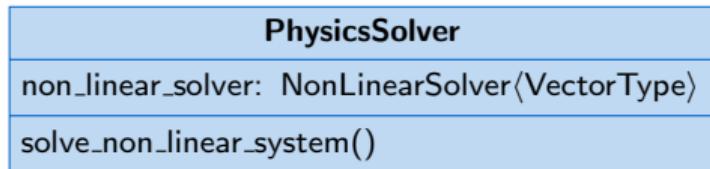
- ▶ We use the deal.II implementation by Kronbichler and Kormann (2012).  
Other libraries: MFEM (Anderson et al., 2021), libCEED (Brown et al., 2021).
  - ▶ Uses cell-quadrature-approach, vectorization and sum-factorization.
- how did we incorporate this idea in Lethe?

## 2. Implementation: from a matrix-based to a matrix-free approach



- ✓ Add new class that uses a system operator instead of a matrix.
- ✓ Mesh and post-processing capabilities already available due to parent class.
- ✓ Several tests from the matrix-based application that can be used to validate and verify the software.

## 2. Implementation: from a matrix-based to a matrix-free approach



- ✗ Lethe supports only Trilinos vectors and the MF infrastructure requires deal.II vectors.
  - ✗ Certain MF functions do not allow the use of non-equal order elements for the velocity and the pressure  
(`MatrixFreeTools::compute_diagonal(...)`).
  - ✗ No preconditioner that supports the matrix-free idea.
- how to implement such a preconditioner?

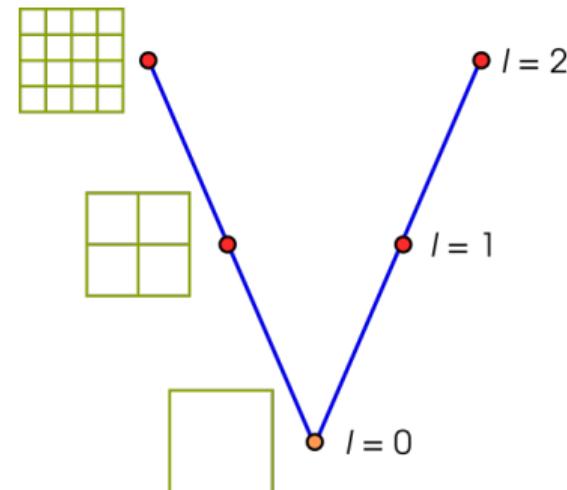
### 3. Linear solver: geometric multigrid preconditioner

A **geometric multigrid** was incorporated in our solver using the implementation available in deal.II (Munch et al., 2023; Clevenger et al., 2020).

**Idea:** correct all components of the errors of the solution using a hierarchy of discretizations based on the geometrical refinement of the grid.

#### Main components:

- ▶ Operators for each level:  $\mathcal{A}^{(l)}, \mathcal{A}^{(0)}$
- ▶ **Smoker** for each level
- ▶ Restriction operator from fine  $l$  to coarse level  $l - 1$ :  $R_l^{(l-1)}$
- ▶ Prolongation operator from coarse  $l - 1$  to fine level  $l$ :  $P_{(l-1)}^l$
- ▶ **Coarse-grid solver**

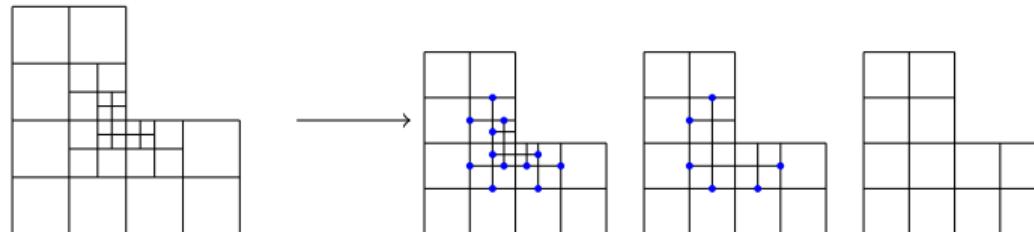


→ for locally-refined meshes there are different geometric multigrid methods

### 3. Linear solver: geometric multigrid preconditioner (Cont.)

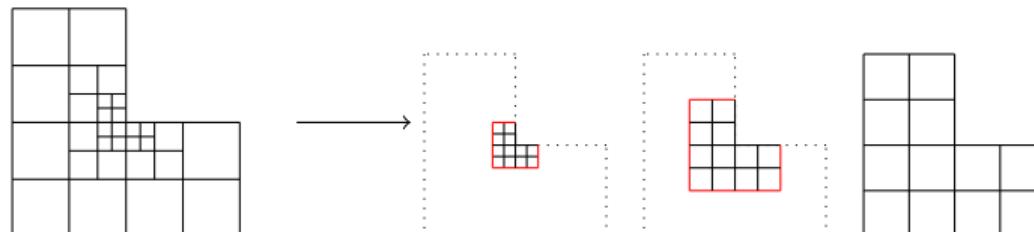
In this work we consider two approaches already implemented in deal.II:

- ▶ **global coarsening (GC):**



No internal interfaces, smoothers consider **hanging nodes**.

- ▶ **local smoothing (LS):**



Internal interfaces (**edges**), efficient transfer between MG levels, no hanging nodes.

→ what do we know about these two approaches?

### 3. Linear solver: geometric multigrid preconditioner (Cont.)

Munch P., Heister T., Prieto Saavedra L. and Kronbichler M. (2023). Efficient distributed matrix-free multigrid methods for locally refined meshes for FEM computations. *ACM Transactions on Parallel Computing*.

- ▶ Considered a Poisson and a Stokes problem with a block preconditioner.
- ▶ Evaluated the performance of both multigrid approaches using different metrics, among them:
  - **Vertical efficiency:** indication of how much data needs to be exchanged during inter-grid transfer → low number: large volume of communication.
  - **Workload efficiency:** measure that considers the amount of cells processed by each processor → low value: load imbalance during smoothing.

#### Main findings:

- ▶ For serial simulations, **LS is faster than GC**: total number of cells to perform smoothing on is lower and there is no cost for evaluating hanging-nodes constraints.
  - ▶ For parallel simulations, **GC is faster than LS**: better load balance.
- **how did we implement both of these approaches in Lethe?**

### 3. Linear solver: geometric multigrid preconditioner (Cont.)

#### NavierStokesPreconditionGMG

mg\_operators: OperatorType  
mg\_smoothen: MGSmoothenPrecondition  
mg\_transfer\_ls: LSTransferType  
mg\_transfer\_gc: GCTransferType  
mg\_coarse: MGCoarseGridBase  
mg: Multigrid  
ls\_mg\_preconditioner: PreconditionMG  
gc\_mg\_preconditioner: PreconditionMG

NavierStokesPreconditionGMG()  
initialize(...)  
vmult(...)

- ▶ In the constructor: we set up the triangulations, the level operators (with constraints), and the transfer operators.
- ▶ In the initialize: transfer relevant quantities to mg levels, evaluate non linear terms for all operators, create smoother (with ev estimation), setup coarse grid solver and create the final preconditioner.

→ This architecture allow us to reuse several parts of the multigrid across transient iterations...

### 3. Linear solver: geometric multigrid preconditioner (Cont.)

#### How to choose the smoother?

- ▶ We use a relaxation scheme  $x_{n+1} = x_n + \omega P^{-1} r_n$  (PreconditionRelaxation), where  $P$  is the diagonal of the operator.
- ▶ We calculate the  $\omega$  using the eigenvalues for every level.

#### What about the coarse-grid solver?

- ▶ Implement different options: AMG, ILU, GMRES preconditioned by AMG/ILU, direct solver.

#### How to further reduce the cost of the preconditioner?

- ▶ Reuse the preconditioner across Newton iterations.
- ▶ Neglect terms involving the Hessian in the jacobian.
- ▶ Use of FE\_Q\_iso\_Q1 elements in coarse-grid level.

→ the solver is tested for different steady/transient problems

## 4. Numerical tests: Taylor-Green vortex

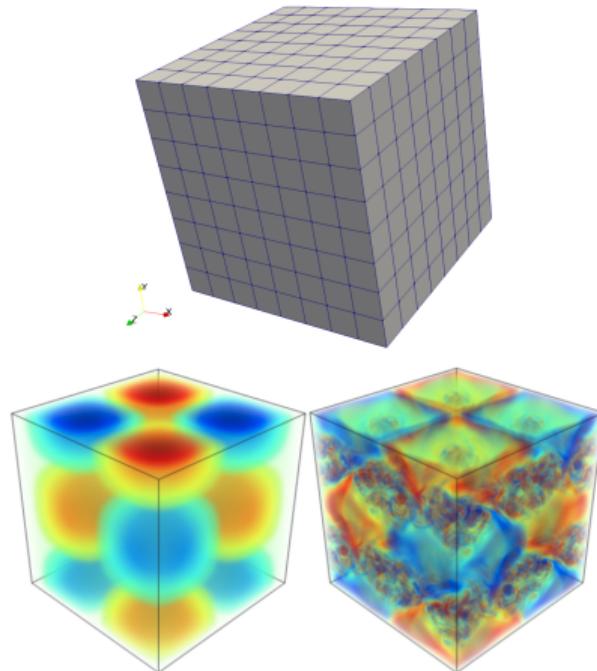
Benchmark to study vortex dynamics and the energy cascade:

- ▶ Transient problem: fixed CFL= 1, BDF2.
- ▶ Domain:  $\Omega = (-\pi, \pi)^3$ .
- ▶ Periodic boundary conditions.
- ▶ Re = 1600.
- ▶  $Q_p Q_p$  elements with  $p = 1, 2, 3$ .
- ▶ Initial condition:

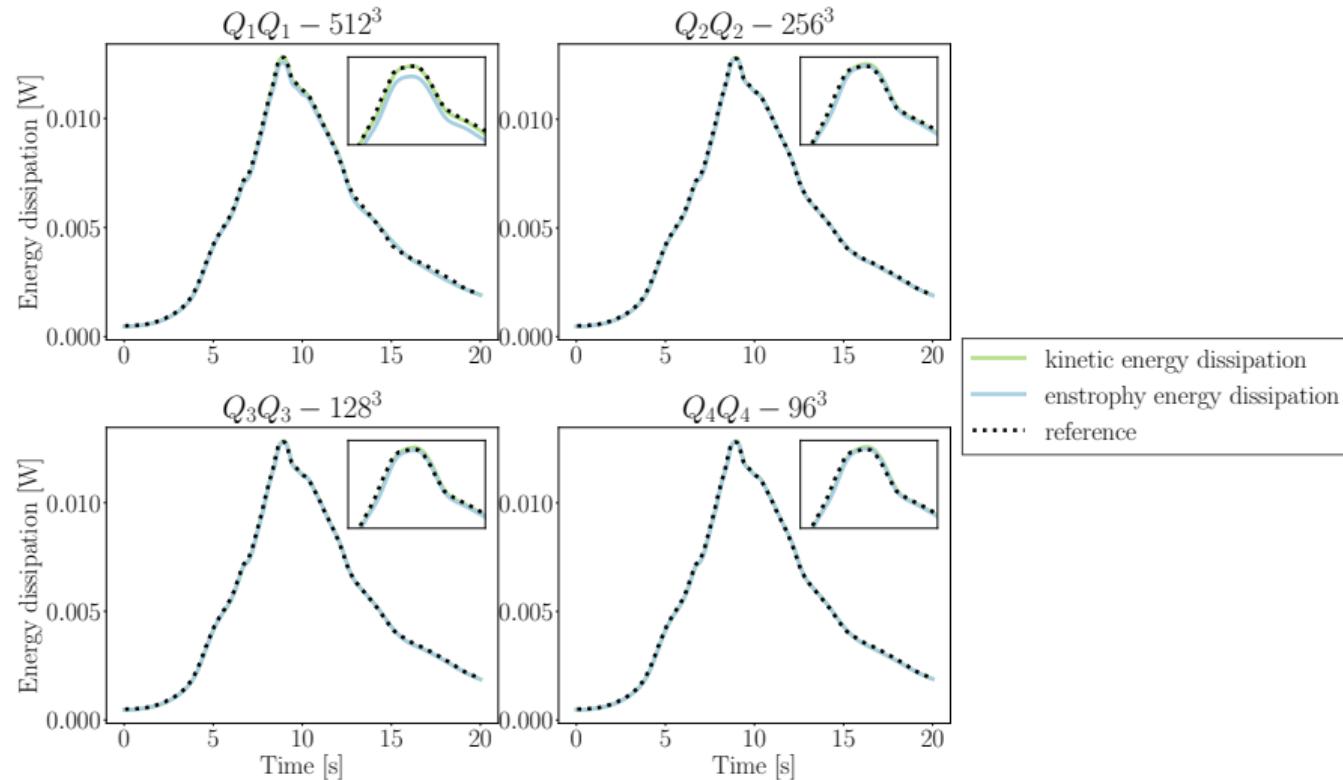
$$\mathbf{u} = \begin{pmatrix} \sin(x) \cos(y) \cos(z) \\ -\cos(x) \sin(y) \cos(z) \\ 0 \end{pmatrix}$$

$$p = \frac{1}{16} (\cos(2x) + \cos(2y)) (\cos(2z) + 2)$$

→ the problem allows us to study parallel scalability

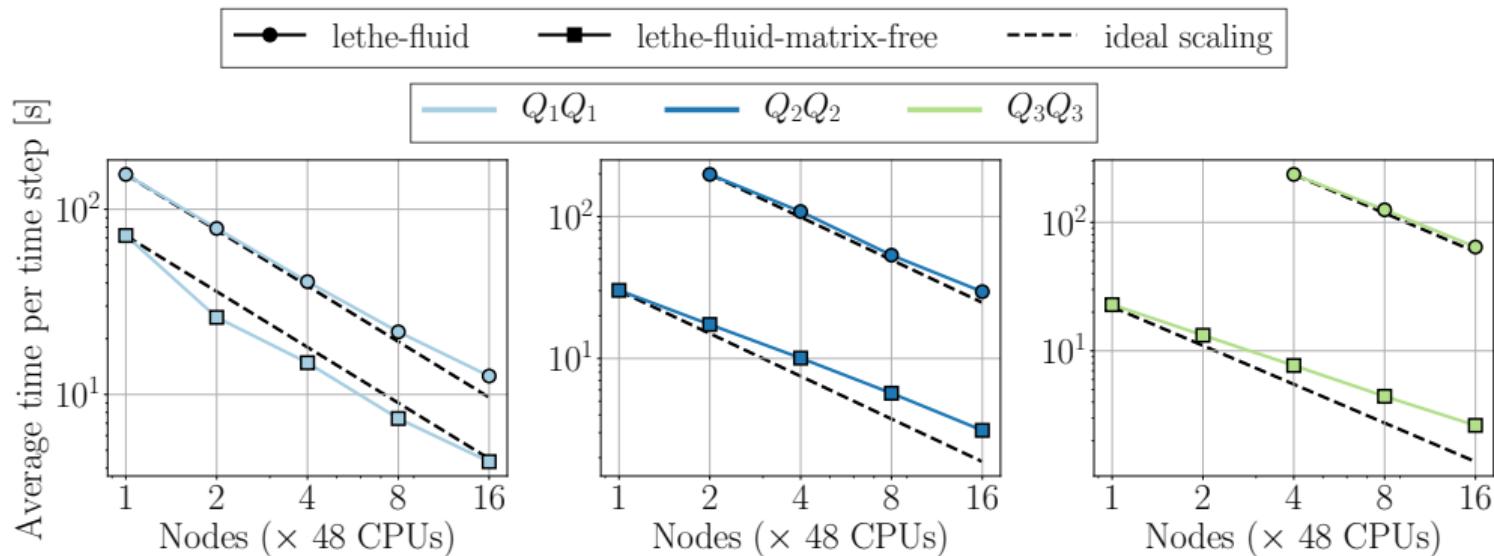


#### 4. Numerical tests: Taylor-Green vortex (Cont.)



- ▶ For the next analysis, we focus on the measurements over the interval  $(6, 10]$ .

#### 4. Numerical tests: Taylor-Green vortex (Cont.)

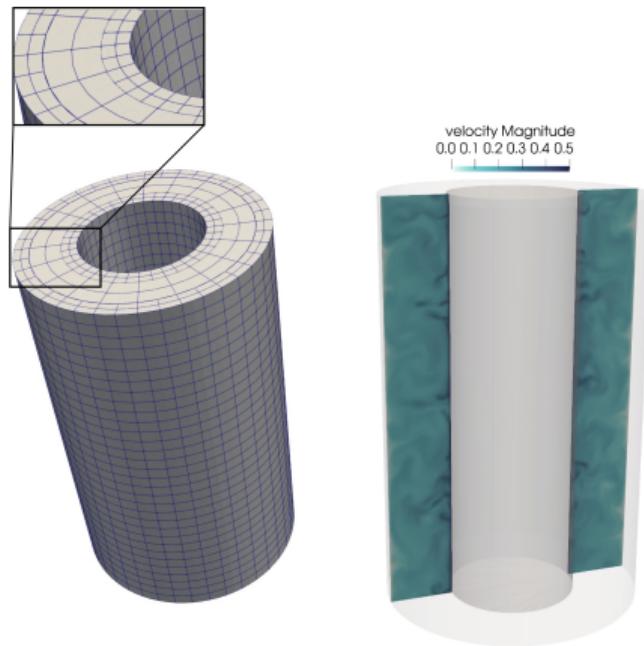


- Matrix-free faster than matrix-based:  $\sim 3x$  for  $Q_1Q_1$ ,  $\sim 10x$  for  $Q_2Q_2$ ,  $\sim 28x$   $Q_3Q_3$ .
- $\sim 34M$  DoFs

## 4. Numerical tests: turbulent Taylor-Couette

Complex turbulent flow problem:

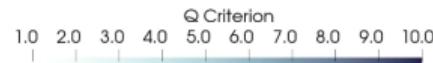
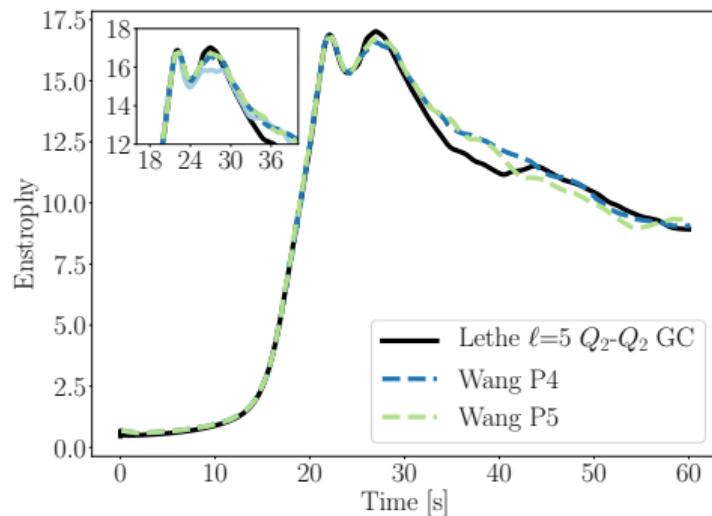
- ▶ Annular flow between two coaxial cylinders.
- ▶ Inner cylinder with a fixed angular velocity.
- ▶ Outer cylinder is static.
- ▶ Curved walls.
- ▶ Transient: BDF2, fixed CFL=1.
- ▶  $\text{Re} = 4000$ .
- ▶  $Q_p Q_p$  elements with  $p = 1, 2$ .
- ▶ **Global static mesh refinement**  $\ell$  with one additional refinement next to the walls.
- ▶ Simulation time: 60s.



→ allow us to study serial and large-scale parallel behavior of the preconditioners

## 4. Numerical tests: turbulent Taylor-Couette - validation

The quality of the results is evaluated through the enstrophy and compared to the reference results from Wang and Jourdan (2021).

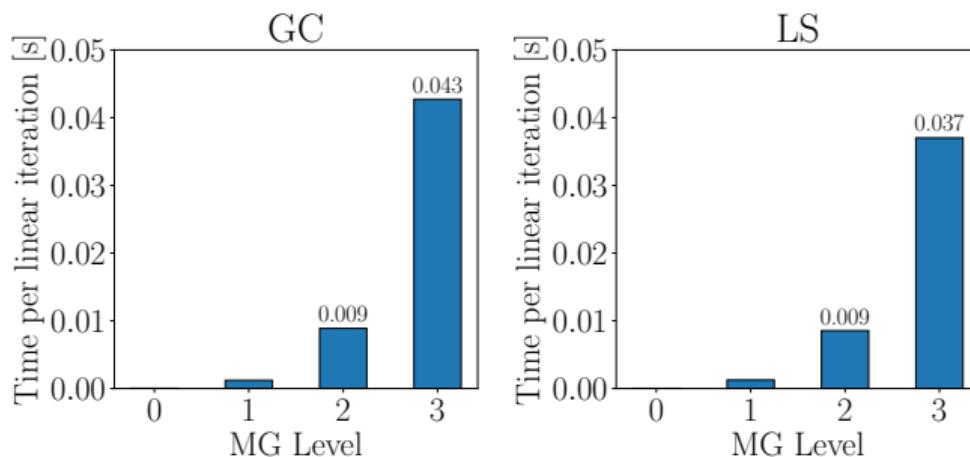


- ▶ For the next analysis, we focus on the measurements over the interval  $(15, 30]$ .

## 4. Numerical tests: turbulent Taylor-Couette - serial run

We analyze average number of linear iterations per Newton iteration (#its) and the total simulation time (t):

$\ell$	$Q_1 Q_1$				$Q_2 Q_2$			
	GC		LS		GC		LS	
	#its	t[s]	#its	t[s]	#its	t[s]	#its	t[s]
2	10	82	15	98.7	9	754	12	640

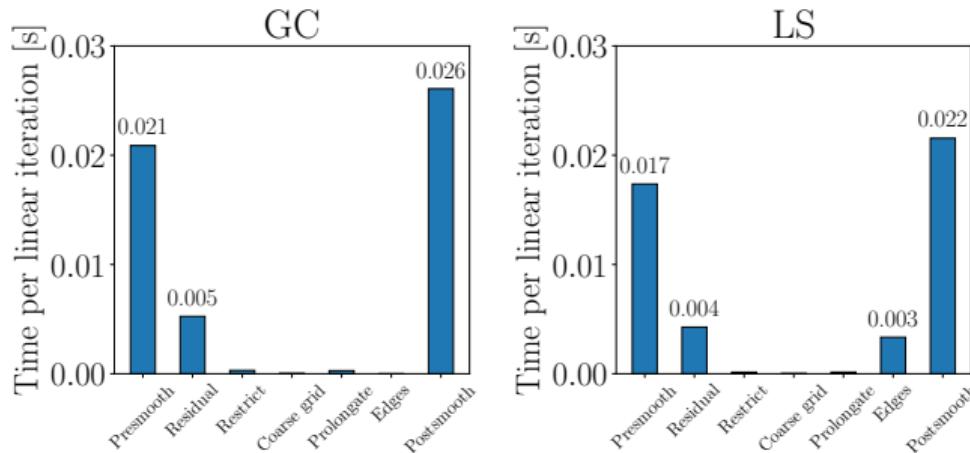


- ▶ GC is faster in terms of total run time (less iterations); but each linear iteration is more expensive.
- ▶ One GC V-cycle is 22% slower than a LS V-cycle:
  - More cells in the case of GC (~12%).
  - Application of expensive hanging node constraints (Munch et al., 2022).

## 4. Numerical tests: turbulent Taylor-Couette - serial run

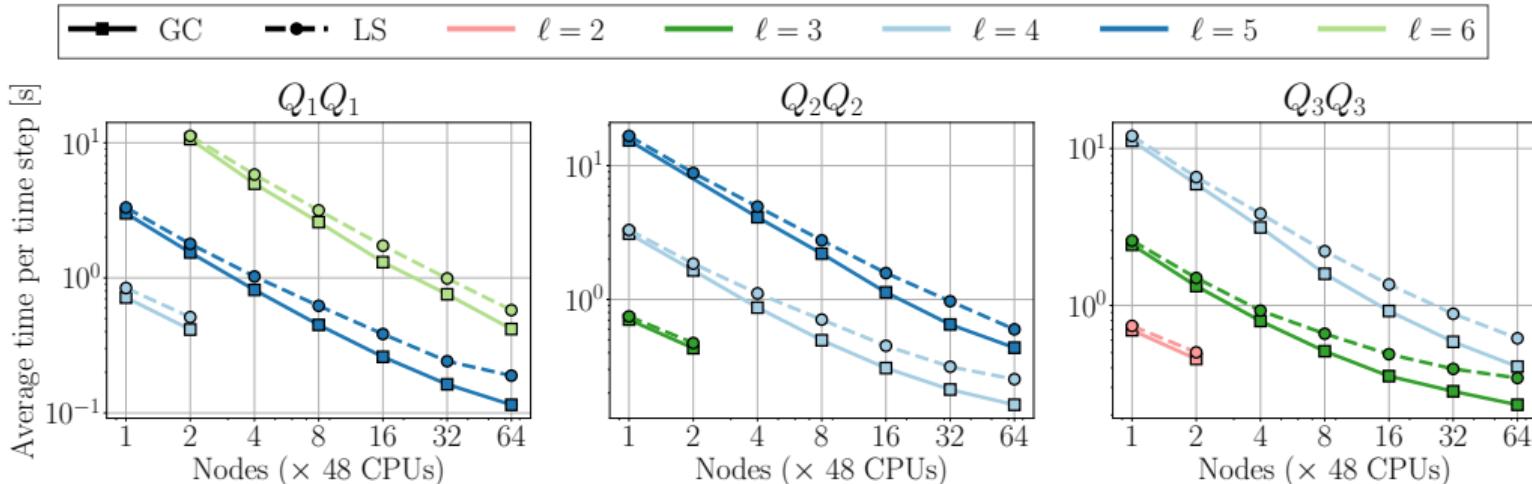
We analyze average number of linear iterations per Newton iteration (#its) and the total simulation time (t):

$\ell$	$Q_1 Q_1$				$Q_2 Q_2$			
	GC		LS		GC		LS	
	#its	t[s]	#its	t[s]	#its	t[s]	#its	t[s]
2	10	82	15	98.7	9	754	12	640



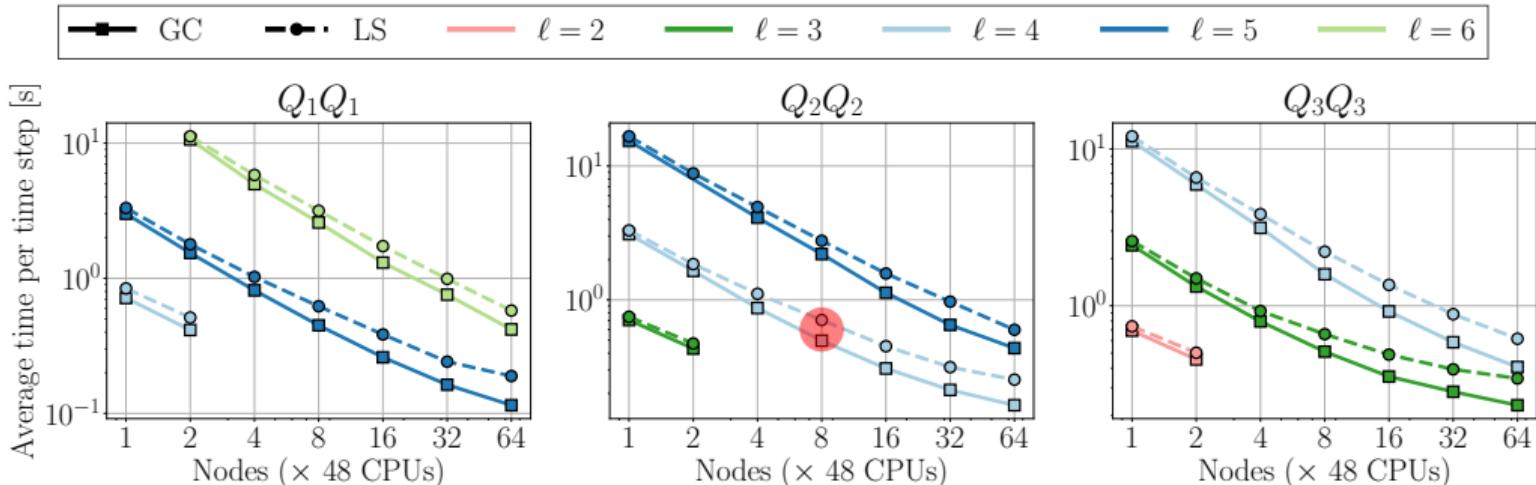
- ▶ Coarse grid is cheap for this case, as it has only 20 cells. Coarse-grid solver: AMG.
- ▶ Pre- and postsmoothing are the most expensive components of the algorithm. LS is faster than GC.
- ▶ Edges (application of non-zero Dirichlet BCs) cost is significant in LS.

#### 4. Numerical tests: turbulent Taylor-Couette - large-scale parallel behavior



- ▶ Up to 6M cells in  $Q_1Q_1$ , 1M cells in  $Q_2Q_2$ , 150K cells in  $Q_3Q_3$
- ▶ GC faster than LS in all cases for all refinements  $\ell$ .
- ▶ GC scales better as we increase the number of nodes.

#### 4. Numerical tests: turbulent Taylor-Couette - large-scale parallel behavior

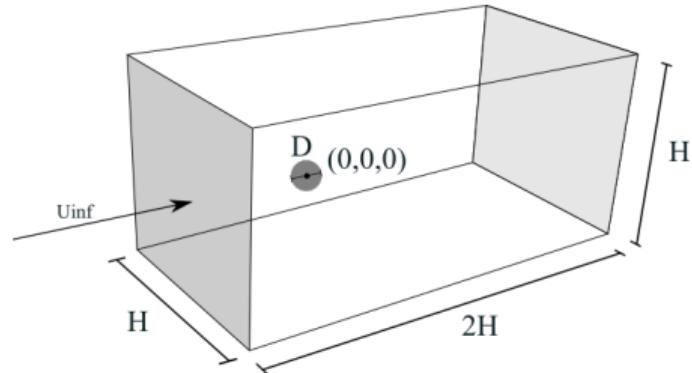
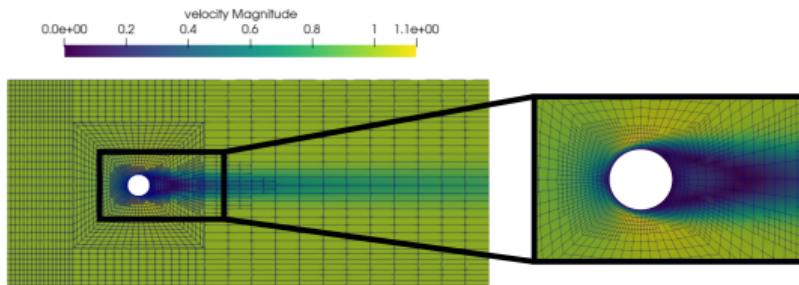


- ▶ Up to 6M cells in  $Q_1Q_1$ , 1M cells in  $Q_2Q_2$ , 150K cells in  $Q_3Q_3$
- ▶ GC faster than LS in all cases for all refinements  $\ell$ .
- ▶ GC scales better as we increase the number of nodes.
- ▶ GC is 1.3x faster than LS. For the same case with one node, LS and GC need the same amount of time.

## 4. Numerical tests: flow around a sphere

Challenging external flow problem:

- ▶ Fixed entrance velocity  $U_\infty = 1$  in the  $x$  direction.
- ▶ Steady-state,  $\text{Re} = 100$ . Initial condition: ramp up  $\text{Re}$  starting with  $\text{Re} = 10$ .
- ▶ No slip boundary conditions around the sphere and slip boundary conditions on the wall.
- ▶  $Q_p Q_p$  elements with  $p = 1, 2$ .
- ▶ Initial global refinement  $\ell$ .
- ▶ **Dynamic mesh refinement:** using Kelly error estimator on the pressure.

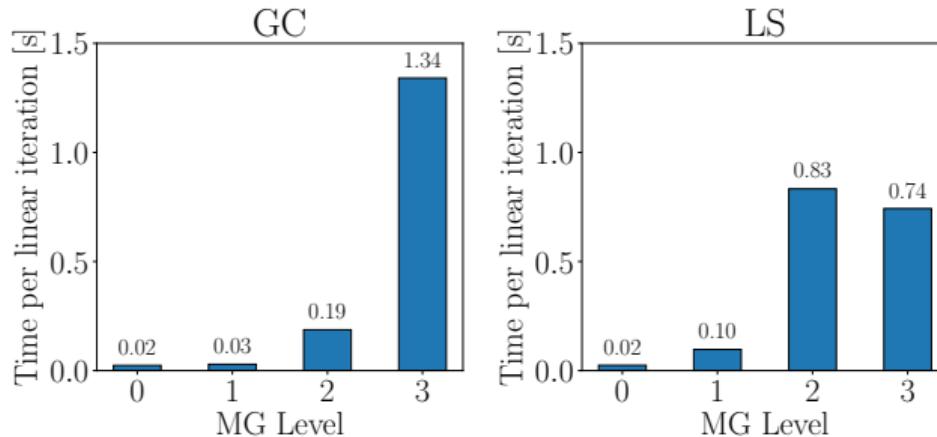


→ allow us to study serial and moderately parallel behavior of the preconditioners

#### 4. Numerical tests: flow around a sphere - serial run

We analyze average number of linear iterations per Newton iteration (#its) and the total simulation time (t):

$\ell$	$Q_1 Q_1$				$Q_2 Q_2$			
	GC		LS		GC		LS	
	#its	t[s]	#its	t[s]	#its	t[s]	#its	t[s]
1	20	37	20	37	24	207	24	207
2	27	173	27	182	42	363	38	346

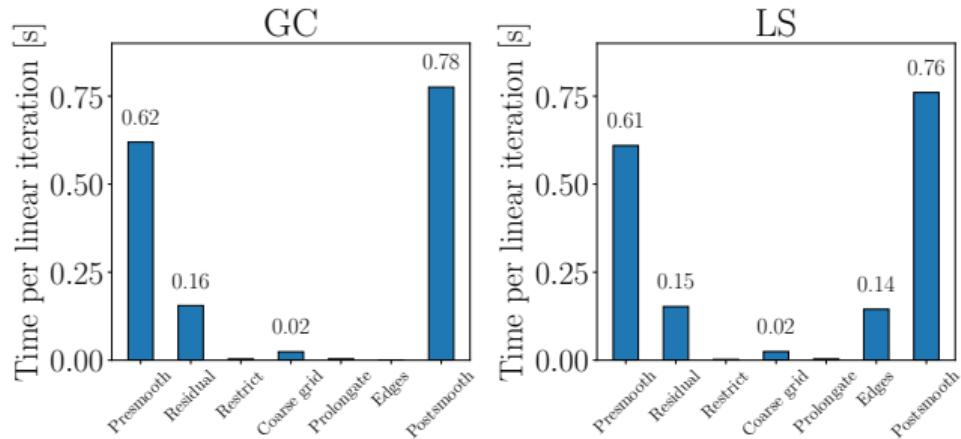


- ▶ In this case, the coarse mesh has 1024 cells. We use GMRES preconditioned by ILU.
- ▶ Approximately same number of iterations and similar time to solution for both methods.
- ▶ Iterations increase a lot with increasing order → possibly due to simple smoother.

## 4. Numerical tests: flow around a sphere - serial run

We analyze average number of linear iterations per Newton iteration (#its) and the total simulation time (t):

$\ell$	$Q_1 Q_1$				$Q_2 Q_2$			
	GC		LS		GC		LS	
	#its	$t[s]$	#its	$t[s]$	#its	$t[s]$	#its	$t[s]$
1	20	37	20	37	24	207	24	207
2	27	173	27	182	42	363	38	346

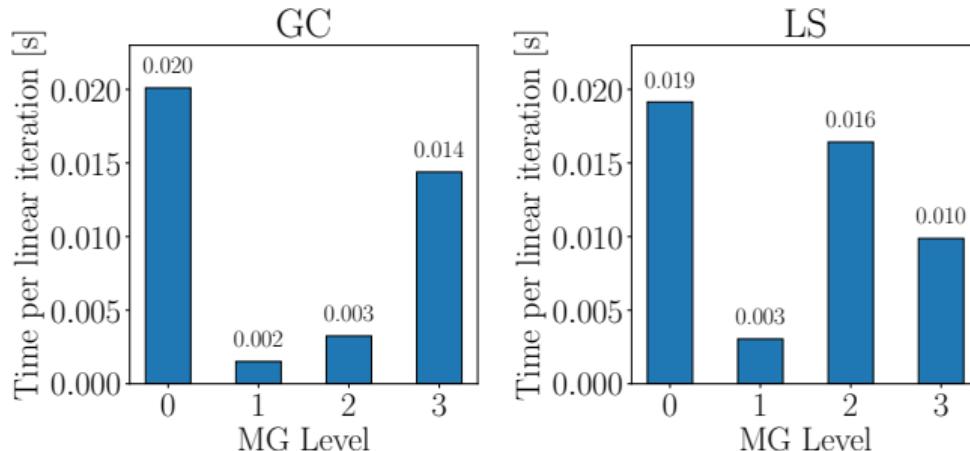


- ▶ Pre- and postsmothing are the most expensive components of the algorithm.
- ▶ Edges cost is significant for LS.

#### 4. Numerical tests: flow around a sphere - parallel run (192 processes)

We analyze average number of linear iterations per Newton iteration (#its) and the total simulation time (t):

$\ell$	$Q_1 Q_1$				$Q_2 Q_2$			
	GC		LS		GC		LS	
	#its	t[s]	#its	t[s]	#its	t[s]	#its	t[s]
1	20	3.46	20	3.46	25	7.84	25	12.7
2	27	4.13	27	5.22	42	7.51	39	12.2



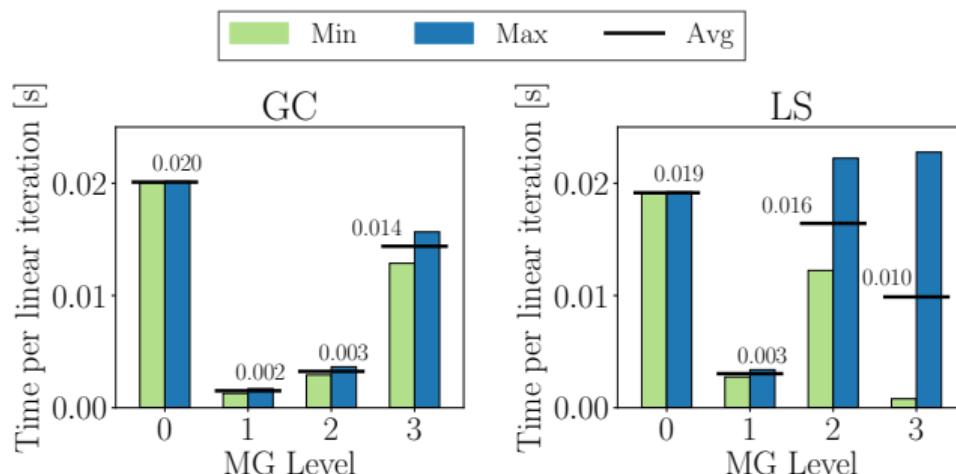
- ▶ Speed up (LS/GC) of 1.2x for  $Q_1 Q_1$ , 1.6x for  $Q_2 Q_2$
- ▶ Same number of iterations for both, but GC faster than LS in all cases.
- ▶ The coarse-grid problem is the bottleneck.

#### 4. Numerical tests: flow around a sphere - parallel run (192 processes)

We analyze average number of linear iterations per Newton iteration (#its) and the total simulation time (t):

$\ell$	$Q_1 Q_1$				$Q_2 Q_2$			
	GC		LS		GC		LS	
	#its	t[s]	#its	t[s]	#its	t[s]	#its	t[s]
1	20	3.46	20	3.46	25	7.84	25	12.7
2	27	4.13	27	5.22	42	7.51	39	12.2

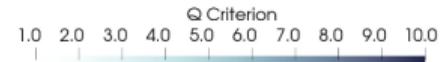
- ▶ LS has load imbalance for the finest levels.
- ▶ GC:
  - Vertical efficiency: 7.48%
  - Workload efficiency: 99%
- ▶ LS:
  - Vertical efficiency: 99.83%
  - Workload efficiency: 52%.
- ▶ The workload efficiency is the most relevant metric, since the smoother is dominating the overall cost of the v-cycle.



## 5. Outlook

### What comes next?

- ▶ Evaluate other smoothers to improve performance for steady-state cases.
- ▶ Evaluate the efficiency of the preconditioners for other benchmarks.
- ▶ Experiment with p-multigrid.
- ▶ Compare stabilization techniques.
- ▶ Couple the matrix-free solver with other physics (e.g., heat transfer, VOF).



Thank you!

Le<sup>o</sup>the<sup>n</sup>

<https://github.com/chaos-polymtl/lethe>

## References I

- Anderson, R., Andrej, J., Barker, A., Bramwell, J., Camier, J.-S., Cerveny, J., Dobrev, V., Dudouit, Y., Fisher, A., Kolev, T., Pazner, W., Stowell, M., Tomov, V., Akkerman, I., Dahm, J., Medina, D., and Zampini, S. (2021). Mfem: A modular finite element methods library. *Computers & Mathematics with Applications*, 81:42–74. Development and Application of Open-source Software for Problems with Numerical PDEs.
- Arndt, D., Bangerth, W., Bergbauer, M., Feder, M., Fehling, M., Heinz, J., Heister, T., Heltai, L., Kronbichler, M., Maier, M., Munch, P., Pelteret, J.-P., Turcksin, B., Wells, D., and Zampini, S. (2023). The deal.II library, version 9.5. *Journal of Numerical Mathematics*, 31(3):231–246.
- Blais, B., Barbeau, L., Bibeau, V., Gauvin, S., El Geitani, T., Golshan, S., Kamble, R., Mirakhori, G., and Chaouki, J. (2020). Lethe: An open-source parallel high-order adaptative cfd solver for incompressible flows. *SoftwareX*, 12:100579.
- Brown, J., Abdelfattah, A., Barra, V., Beams, N., Camier, J.-S., Dobrev, V., Dudouit, Y., Ghaffari, L., Kolev, T., Medina, D., Pazner, W., Ratnayaka, T., Thompson, J., and Tomov, S. (2021). libCEED: Fast algebra for high-order element-based discretizations. *Journal of Open Source Software*, 6(63):2945.
- Clevenger, T. C., Heister, T., Kanschat, G., and Kronbichler, M. (2020). A flexible, parallel, adaptive geometric multigrid method for fem. *ACM Trans. Math. Softw.*, 47(1).
- Kronbichler, M. and Kormann, K. (2012). A generic interface for parallel cell-based finite element operator application. *Computers and Fluids*, 63:135–147.
- Munch, P., Heister, T., Saavedra, L. P., and Kronbichler, M. (2023). Efficient distributed matrix-free multigrid methods on locally refined meshes for FEM computations. *ACM Transactions on Parallel Computing*, pages 1–34.
- Munch, P., Ljungkvist, K., and Kronbichler, M. (2022). Efficient application of hanging-node constraints for matrix-free high-order fem computations on cpu and gpu. In *High Performance Computing: 37th International Conference, ISC High Performance 2022, Hamburg, Germany, May 29 – June 2, 2022, Proceedings*, page 133–152, Berlin, Heidelberg. Springer-Verlag.
- Wang, Z. J. and Jourdan, E. (2021). Benchmark for scale-resolving simulation with curved walls: the Taylor Couette flow. *Advances in Aerodynamics*, 3(1).

## Numerical implementation: the cell matrix

In reality, the local element matrix can be decomposed as the product of three matrices:

$$\mathbf{A}_k = \mathcal{S}_K^T \mathcal{Q}_K \mathcal{S}_K \quad (1)$$

But how? For this, we need to take a closer look to the diffusion term in the Jacobian expression:

$$F'(\mathbf{u}, p)[\delta \mathbf{u}, \delta p] = \cdots + \nu(\nabla \mathbf{v}, \nabla \delta \mathbf{u})_{\Omega} + \cdots \quad (2)$$

If we insert the discretization, this looks like this:

$$\begin{aligned} \nu(\nabla \mathbf{v}, \nabla \delta \mathbf{u})_{\Omega} &= \nu \sum_k (\nabla_x \mathbf{v}_k, \nabla_x \mathbf{u}_k)_{\Omega_k} \\ &= \nu \sum_k V_{k,i}^T (\nabla_x \phi_i, U_{k,j} \nabla_x \phi_j)_{\Omega_k} \\ &= \nu \sum_k V_{k,i}^T \underbrace{\int_{\Omega_k} (\nabla_x \phi_i)^T \nabla_x \phi_j U_{k,j} dx}_1 \end{aligned}$$

where  $\nabla_x$  is the gradient in physical space.

→ now we use the **cell-quadrature approach** for 1

## 4. Numerical implementation: cell-quadrature approach

**Reminder:** in finite elements we define a **mapping** to go from the physical space  $x$  to the reference space  $\xi$ , which introduces the Jacobian matrix  $J$  in several terms. In addition, we use **numerical quadrature** to approximate the integral:  $\int F(\xi) d\xi \simeq \sum_q \omega_q g(\xi_q)$ .

We can further develop the expression as:

$$\begin{aligned} \int_{\Omega_k} (\nabla_x \phi_i)^T \nabla_x \phi_j U_{k,j} dx &= \int_{\Omega_k} \left( J_k^T \nabla_\xi \phi_i \right)^T \left( J_k^{-T} \nabla_\xi \phi_j \right) U_{k,j} (\det J_k) d\xi \\ &= \sum_q \underbrace{(\nabla_{\xi_q} \phi_i)^T}_{S_k^T} \underbrace{J_k^T \omega_q (\det J_k) J_k^{-T}}_{Q_k} \underbrace{(\nabla_{\xi_q} \phi_j)}_{S_k} U_{k,j} \end{aligned}$$

From left to right:

1. Transform the vector values on the local DoFs to a vector of gradients in quadrature points.
2. Multiply the gradients by the integration weights and Jacobian information.
3. Apply the second gradient to have a vector of Laplacian values on the cell DoFs.

## Multigrid algorithm

---

**Algorithm 1** Solves  $Ax = b$  approximately

---

```
1: function MultigridVCycle( $I$ ,  $A^{(I)}$ ,  $x^{(I)}$ ,  $b^{(I)}$ )
2: if  $I = 0$  then
3:    $x^{(0)} \leftarrow \text{CoarseGridSolver}(A^{(0)}, x^{(0)}, b^{(0)})$  (e.g. Algebraic Multigrid (AMG))
4: else
5:    $x^{(I)} \leftarrow \text{Smoothen}(A^{(I)}, x^{(I)}, b^{(I)})$ 
6:    $r^{(I)} \leftarrow b^{(I)} - A^{(I)} x^{(I)}$ 
7:    $b^{(I-1)} \leftarrow R_I^{(I-1)} r^{(I)}$ 
8:    $x^{(I-1)} \leftarrow \text{MultigridVCycle}(I-1, A^{(I-1)}, 0, b^{(I-1)})$ 
9:    $x^{(I)} \leftarrow x^{(I)} + P_{(I-1)}^I x^{(I-1)}$ 
10:   $x^{(I)} \leftarrow \text{Smoothen}(A^{(I)}, x^{(I)}, b^{(I)})$ 
11:  return  $x^{(I)}$ 
12: end if
```

---

→ the difference between multigrid methods relies in the way the **hierarchy of discretizations** is defined

## Flow around a sphere: pressure coefficient

$$C_D = C_p + C_\tau = \frac{F_p}{\frac{1}{2}\rho U_\infty^2 (\frac{\pi}{4}D^2)} + \frac{F_\tau}{\frac{1}{2}\rho U_\infty^2 (\frac{\pi}{4}D^2)}$$

■  $Q_1 Q_1$       ♦  $Q_2 Q_2$       ○  $Q_3 Q_3$

